



Deep Learning Basics

Datta Lab presentation @ Princeton University

2023-07-07

Martin Lellep
martin.lellep@gmail.com

Purpose of this talk

Teach you the very basic idea behind
Deep Learning (DL)

Course materials are available!



<http://lellep.xyz/blog/datta-lab-ml-course.html>

Questions

- 1) Who has **heard** of neural networks?
- 2) Who has **used** neural networks in some way?

Why are Neural Networks cool?

- Ever used translator? → Neural network
- Ever Amazon suggestions? → Neural network
- Ever used subtitles on YouTube? → Neural network

Conclusion: Very powerful tools!




They must be useful for Physics, too!

Why are Neural Networks cool?



<https://www.marburg.de/portal/seiten/webcam-blick-aufs-landgrafenschloss-und-den-marktplatz-900001035-23001.html>



Object detection applied to view on Marburg (=city in Germany)

Agenda

- Introduction
- Neural network (NN) structure
- Stochastic gradient descent (SGD)
- Back propagation (backprop)
- Frameworks
- Deep Learning (DL) examples


This scratches only the surface!

Introduction

- Examples of simple function approximators:
 - Linear function: $f(x) = ax + b$
 - Quadratic function: $f(x) = ax^2 + bx + c$
- All of them:
 - Parameterised by a set of parameters
 - Can be used to fit given data, $\{(x, y)_i\}_i$

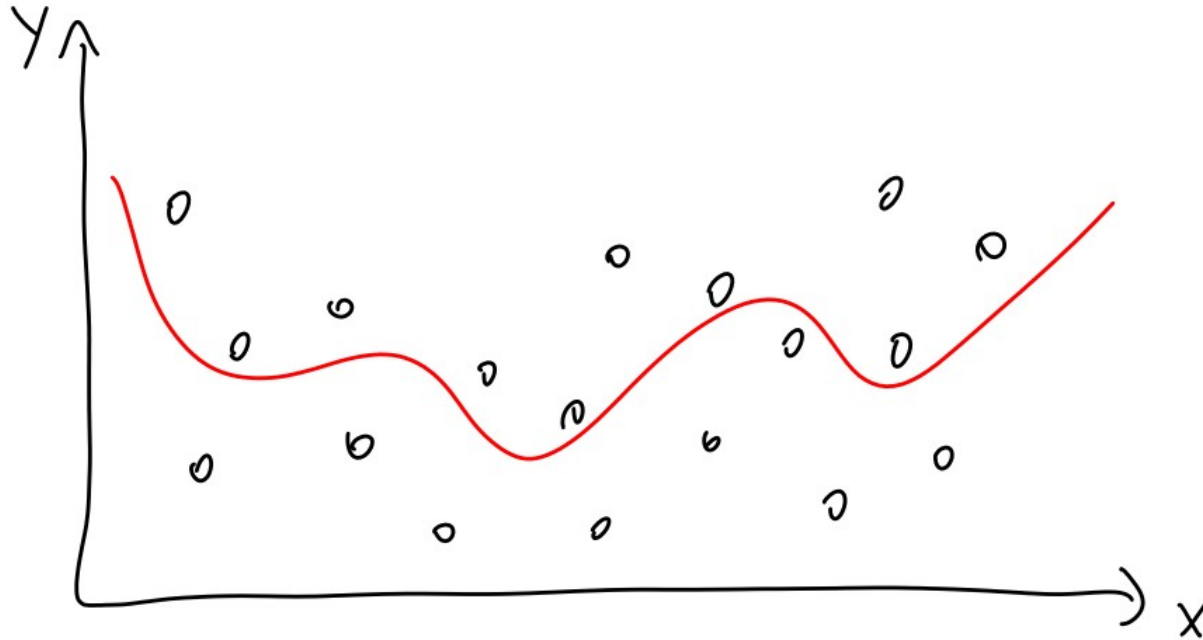
Introduction

Example of fitting a linear function:



Introduction

Example of fitting a non-linear function:



Introduction

[Live demo – Fitting function]

Introduction


- Steps for fitting a function
 - Data to fit
 - A function to fit – called **model** for here
 - An optimisation procedure to perform the fitting
 - A cost function that is minimised
- Question for this course: How to choose model?
- Solution: Use NNs and DL!

NN structure

- A NN is a function approximator
 - ... a complicated function approximator
 - ... a parameterised function approximator
 - ... a non-linear function approximator
- In fact, so complicated that considered a black box for the most part


NN structure

Black box character of NNs:




NN structure

- NN structure inspired by neurons in brain
- NNs consist of layers of neurons:



NN structure

Focus on one single layer:



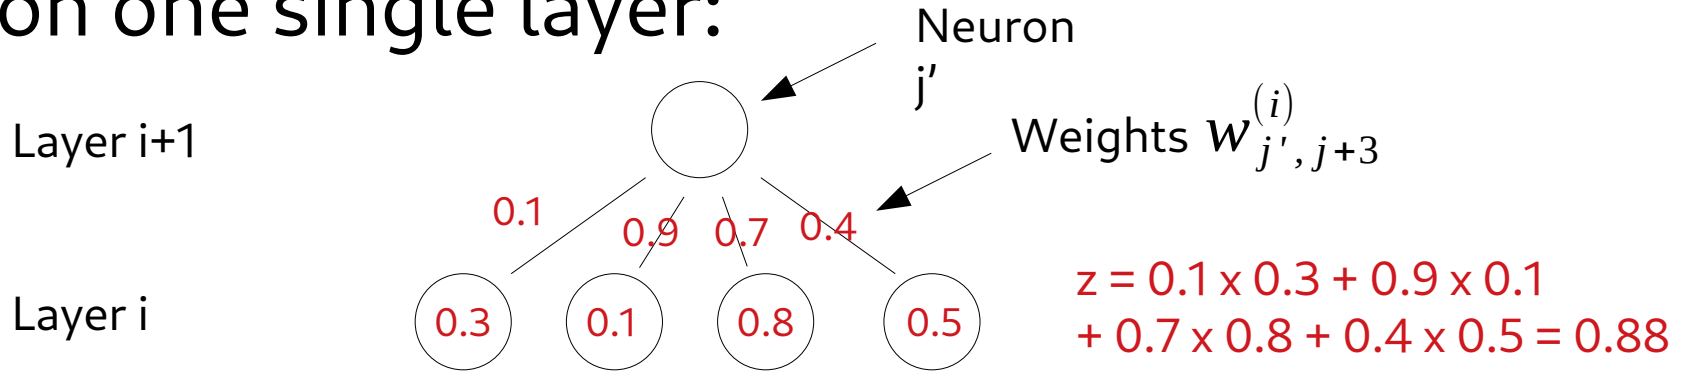
NN structure

(Forward) evaluation of NN:

- 1) Linear combination of previous layer
- 2) Activation function applied to step 1)

NN structure

Focus on one single layer:



Neuron: j $j+1$ $j+2$ $j+3$

Linear combination: $z_{j'}^{(i+1)} = w_{j',j}^{(i)} y_j^{(i)} + w_{j',j+1}^{(i)} y_{j+1}^{(i)} + w_{j',j+2}^{(i)} y_{j+2}^{(i)} + w_{j',j+3}^{(i)} y_{j+3}^{(i)}$


Non-linear function ("activation"): $y_{j'}^{(i+1)} = g(y_{j'}^{(i+1)})$

Note: Almost all linear operations
can be cast into linear algebra
operations => very fast!

NN structure

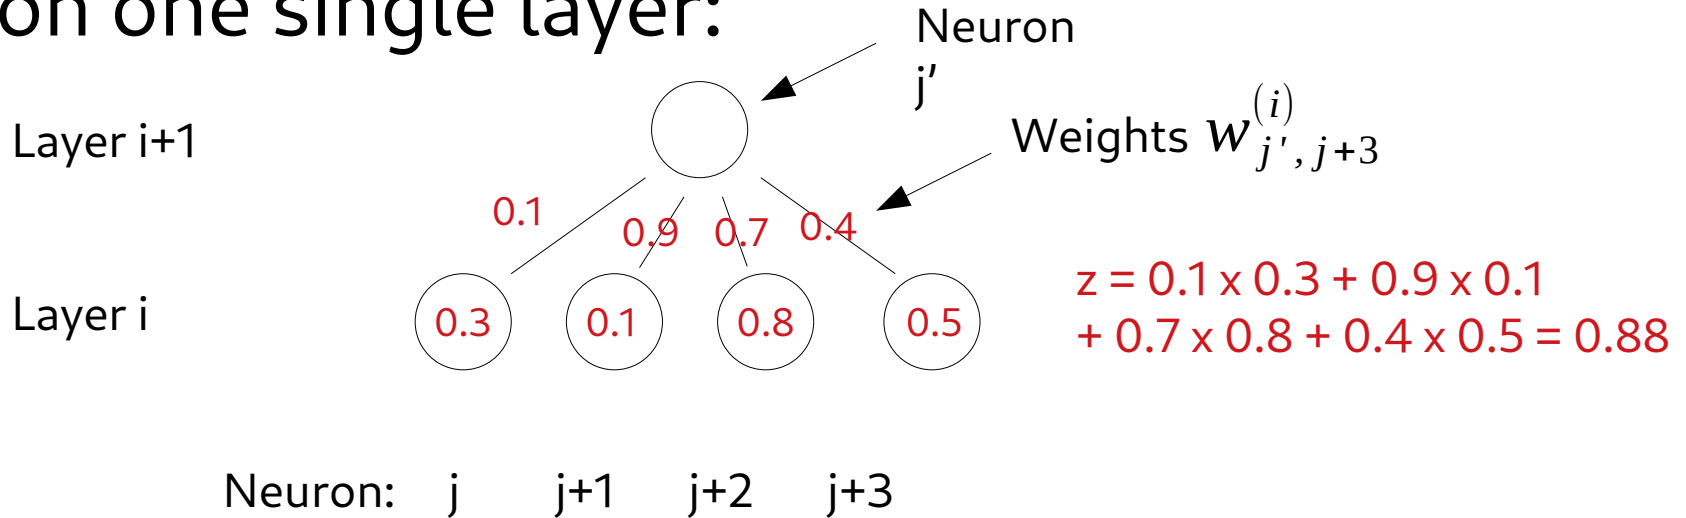
Activation functions:

- Anything close to a step function
- Fast to compute
- Easy to differentiate



NN structure

Focus on one single layer:



Linear combination: $z_{j'}^{(i+1)} = w_{j',j}^{(i)} y_j^{(i)} + w_{j',j+1}^{(i)} y_{j+1}^{(i)} + w_{j',j+2}^{(i)} y_{j+2}^{(i)} + w_{j',j+3}^{(i)} y_{j+3}^{(i)}$

Non-linear function ("activation"): $y_{j'}^{(i+1)} = g(y_{j'}^{(i+1)})$ $y = \text{sigma}(0.88) = 0.71$

NN structure

Intermediate summary:

- NNs are built from layers
- Weights connect adjacent layers
- Information flows from input to output

remember: linear vs
quadratic

Model

Parameters

Remaining: also called "loss"

- Cost function:

$$L = \frac{1}{2} \sum_{i=1}^N (y_i - y_{NN}(x_i; \theta))^2$$

True output
value of sample i

NN prediction for
output value of
sample i

Sum over training samples

Values of all
weight matrices


NN structure

Summary:


- We know how to evaluate a given feed forward fully connected NN

Now:

- Fit NN to regression task



Called "training" in
analogy to
human learning




Classical optimisation
task

Stochastic Gradient Decent (SGD)

- Optimisation procedure to minimise cost function L
- Typical first-order optimisation routine: gradient descent
 - Follow negative gradient
 - Update: $\theta' = \theta - \epsilon \nabla_{\theta} L(\theta)$

Step size or “learning rate”



SGD

- Key problem: $\theta' = \theta - \epsilon \nabla_{\theta} L(\theta)$
- Why?

$$L = \frac{1}{2} \sum_{i=1}^N (y_i - y_{NN}(x_i; \theta))^2$$

← Veeeeeeeery large! $\sim O(1e5)$

- Solution:

batch size

 - Use $m \ll N$ randomly chosen samples

$$L = \frac{1}{2} \sum_{i=1}^m (y_i - y_{NN}(x_i; \theta))^2$$

SGD

Then:

$$\begin{aligned}\nabla_{\theta} L &= \frac{1}{2} \sum_{i=1}^m \nabla_{\theta} (y_i - y_{NN}(x_i; \theta))^2 \\ &= - \sum_{i=1}^m (y_i - y_{NN}(x_i; \theta)) \nabla_{\theta} y_{NN}(x_i; \theta)\end{aligned}$$

The neural network structure is known.

Hence: It is possible to compute it!

SGD

Then:

$$\begin{aligned}\nabla_{\theta} L &= \frac{1}{2} \sum_{i=1}^m \nabla_{\theta} (y_i - y_{NN}(x_i; \theta))^2 \\ &= - \sum_{i=1}^m (y_i - y_{NN}(x_i; \theta)) \nabla_{\theta} y_{NN}(x_i; \theta)\end{aligned}$$

The neural network structure is known.

Hence: It is possible to compute it!

Stunning side note: SGD helps with avoiding to get stuck at saddle points or (weak) local minima.

“Backprop” to the rescue!

Backprop


- “Backprop” = back propagation
- Very rough concept:
 - Forward pass: Feed data into the network
 - Backward pass: Propagate errors backwards to adapt weights
- Remember for now:
 - Backprop makes computing the derivatives easy
 - It is **much simpler** than computing the gradient naively

Frameworks

- Nobody uses plain Python implementations
- Nobody who does NN and DL in 2023 at least
- Solution: e.g. Tensorflow and Keras!

Deals with tensors
and differentiation

Deals with NNs



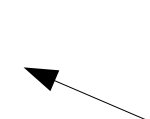


Keras

Frameworks: Keras

- High-level Python library
- Useful for:
 - Assembly of NN
 - Training of NN
 - Evaluation of NN (training vs validation vs test data)
 - And some others

Frameworks: Keras

- NNs, as presented before, are modeled as layers
- The user has to provide data
- Data:
 - Training data  Used for training
 - Validation data  Used for design optimisations of NN using **unseen** data
 - Test data  Test of how well the NN performs on **unseen** data

Frameworks: Keras

[Live demo – Learn 2D function]


Deep learning (DL)

- DL = Use NNs that have many layers
- Previously, we discussed “feed forward” layers (called “Dense” in Keras)
- But: There are many more types of layers
 - For images (ConvLayers)
 - For time series (LSTMs)

In fact: You are only limited by your imagination when designing layers (if they perform well is, of course, a different question)

Convolutional NNs (CNNs)

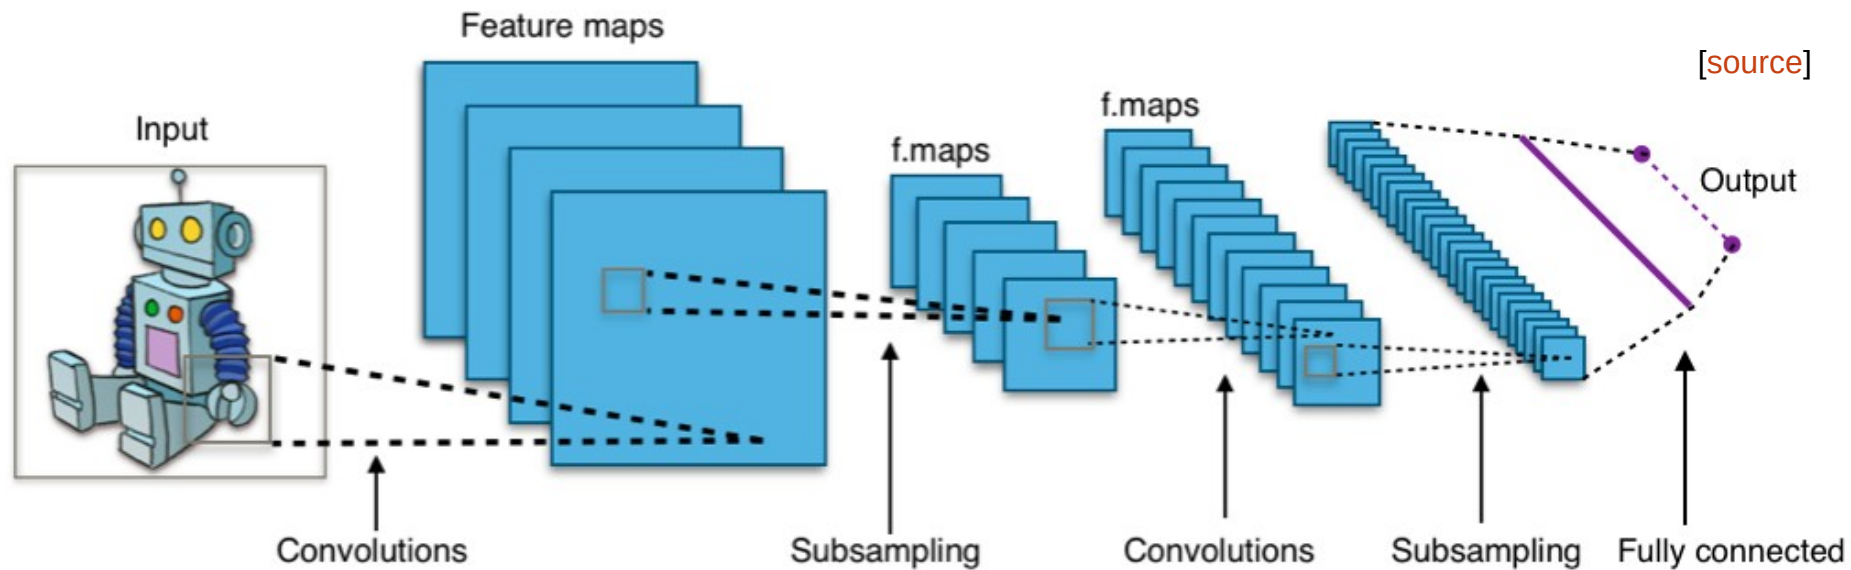
- A layer type that learns convolution kernels
- Typically used for images
- Advantage:
 - Less weights than Dense
 - Good for translationally invariant features



[source]


Another nice visualisation!

CNNs




DL vs ML

Machine Learning




Deep Learning



→ DL learns the features by itself!

MNIST dataset

- A dataset to test machine learning model performances
 - Widely used in DL to benchmark how well a NN performs
 - Content: Images of handwritten digits
- Examples:



Advanced example

- We train a CNN on MNIST data
 - We use different layer types
 - We use other best practises
 - We use a larger model!
- ... and see some drawbacks of NNs



They are computationally expensive to train

Frameworks: Keras

[Live demo – MNIST with CNN]

Colab

- A platform provided by Google
- I presented the examples using **Jupyter Notebooks**
 - “Python in the cloud”
- Colab uses Jupyter Notebooks that run on Google infrastructure
- Advantage: They provide free GPU access whi

<http://colab.research.google.com>

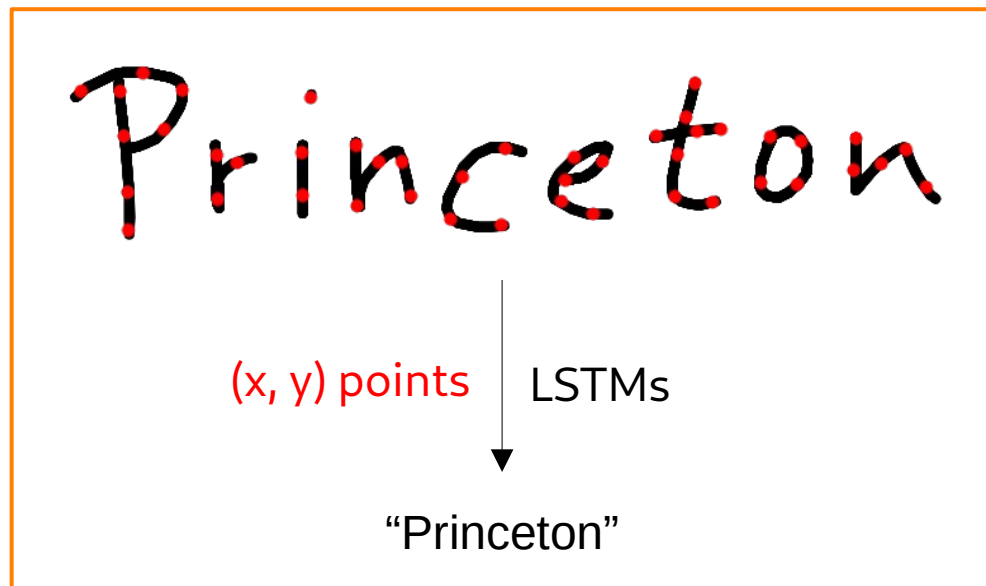
Outlook

- There is much more to DL
- Examples:
 - LSTM layers
 - Autoencoder
 - GANs
 - VAEs
 - Information theory (to understand more details)

Re LSTMs: my weekend project

Small Sunday-only hobby project of mine:

- Given pen dynamics, predict written text
- Uses LSTMs
- Impact: O(10,000) users of open source application



Project link & demo video:

https://github.com/PellelNitram/xournalpp_htr

Sources

- Backprop:
 - <https://mattmazur.com/2015/03/17/a-step-by-step-backpropagation-example/>
- CNNs:
https://de.wikipedia.org/wiki/Convolutional_Neural_Network
- CNN & MNIST example:
 - https://keras.io/examples/mnist_cnn/
 - Colab: <https://colab.research.google.com>

Get in contact

E-Mail: martin.llelep@gmail.com



Tech Blog:
<http://llelep.xyz/blog>



LinkedIn:
<https://www.linkedin.com/in/martin-llelep-858600152/>

Questions



Feel free to ask questions!

Next session

Neural networks
& deep learning